Hadoop File Formats Practical

...

IMPORTANT

Copyright Infringement and Illegal Content Sharing Notice

Orc File Format

Create a table with orc file format named “orders_orc”

```
CREATE TABLE orders_orc(
  id bigint,
  product_id string,
  customer_id bigint,
  quantity int,
  amount double) stored as orc;
```

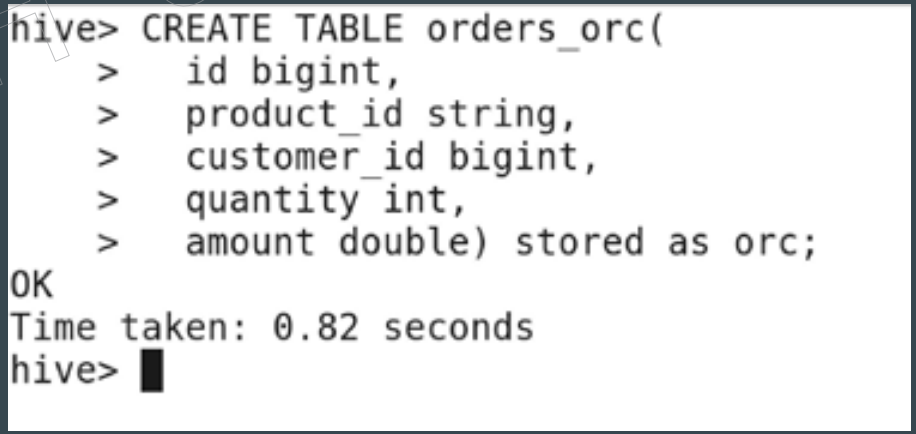

Orc File Format

Now insert the data in this table from orders table

insert into orders_orc select * from orders;

```
hive> insert into orders_orc select * from orders;
Query ID = cloudera_20200507031616_8722e716-a7ec-46a1-9e83-ac1b8221618a
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1588458683533_0011, Tracking URL = http://quickstart.cloudera
88/proxy/application_1588458683533_0011/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1588458683533_0011
```

Orc File Format

Now try to see the data in hdfs

hadoop fs -ls /user/hive/warehouse/trendytech.db/orders_orc/000000_0

hadoop fs -cat /user/hive/warehouse/trendytech.db/orders_orc/*

```
[cloudera@quickstart ~]$ hadoop fs -ls /user/hive/warehouse/trendytech.db/orders_orc
Found 1 items
-rwxrwxrwx   1 cloudera supergroup        645 2020-05-07 03:29 /user/hive/warehouse/trendytech.db/ord
ers_orc/000000_0
[cloudera@quickstart ~]$
```

Orc File Format

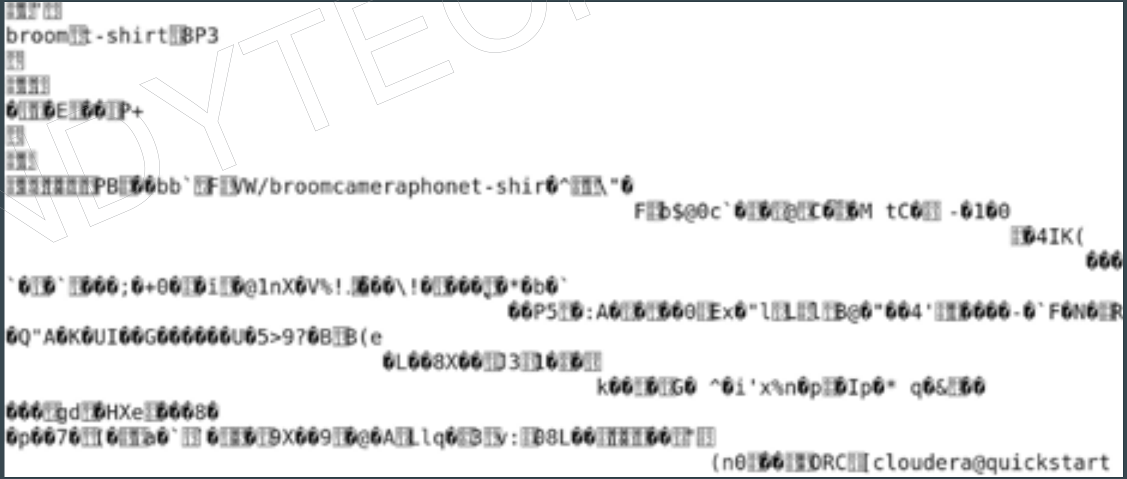

Orc File Format

To get information about an ORC file, use the orcfiledump command

hive --orcfiledump /user/hive/warehouse/trendytech.db/orders_orc/000000_0

```
[cloudera@quickstart ~]$ hive --orcfiledump /user/hive/warehouse/trendytech.db/orders_orc/000000_0
Structure for /user/hive/warehouse/trendytech.db/orders_orc/000000_0
File Version: 0.12 with HIVE_8732
20/05/07 03:38:54 INFO orc.ReaderImpl: Reading ORC rows from /user/hive/warehouse/trendytech.db/orders_orc/000000_0 with {include: null, offset: 0, length: 9223372036854775807}
Rows: 5
Compression: ZLIB
Compression size: 262144
Type: struct<_col0:bigint,_col1:string,_col2:bigint,_col3:int,_col4:double>

Stripe Statistics:
  Stripe 1:
    Column 0: count: 5 hasNull: false
    Column 1: count: 5 hasNull: false min: 111111 max: 111115 sum: 555565
    Column 2: count: 5 hasNull: false min: broom max: t-shirt sum: 28
    Column 3: count: 5 hasNull: false min: 1111 max: 4444 sum: 9999
    Column 4: count: 5 hasNull: false min: 1 max: 3 sum: 9
    Column 5: count: 5 hasNull: false min: 10.0 max: 5200.0 sum: 6496.0

File Statistics:
  Column 0: count: 5 hasNull: false
```

Orc File Format

To display the data in the ORC file, use

hive --orcfiledump -d /user/hive/warehouse/trendytech.db/orders_orc/000000_0

```
[cloudera@quickstart ~]$ hive --orcfiledump -d /user/hive/warehouse/trendytech.db/orders_orc/000000_0
20/05/07 03:47:20 INFO orc.ReaderImpl: Reading ORC rows from /user/hive/warehouse/trendytech.db/order
s_orc/000000_0 with {include: null, offset: 0, length: 9223372036854775807}
{"_col0":111111,"_col1":"phone","_col2":1111,"_col3":3,"_col4":1200}
{"_col0":111112,"_col1":"camera","_col2":1111,"_col3":1,"_col4":5200}
{"_col0":111113,"_col1":"broom","_col2":1111,"_col3":1,"_col4":10}
{"_col0":111114,"_col1":"broom","_col2":2222,"_col3":2,"_col4":20}
{"_col0":111115,"_col1":"t-shirt","_col2":4444,"_col3":2,"_col4":66}
[cloudera@quickstart ~]$
```

Parquet File Format

Create a table with parquet file format named “orders_parquet”

```
CREATE TABLE orders_parquet(
  id bigint,
  product_id string,
  customer_id bigint,
  quantity int,
  amount double) stored as parquet;
```

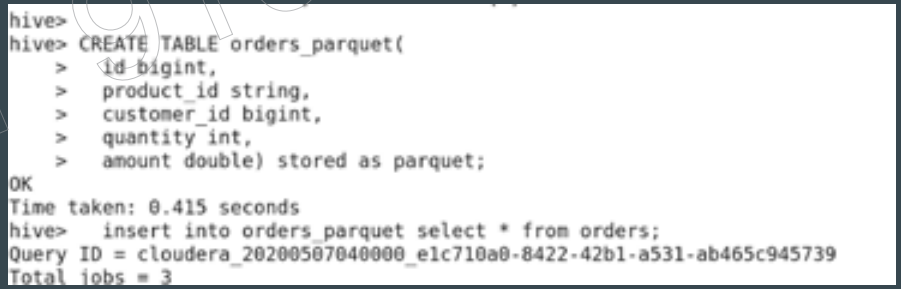


```
insert into orders_parquet select * from orders;
```

Parquet File Format

Now try to see the data in hdfs

hadoop fs -cat /user/hive/warehouse/trendytech.db/orders_parquet/*

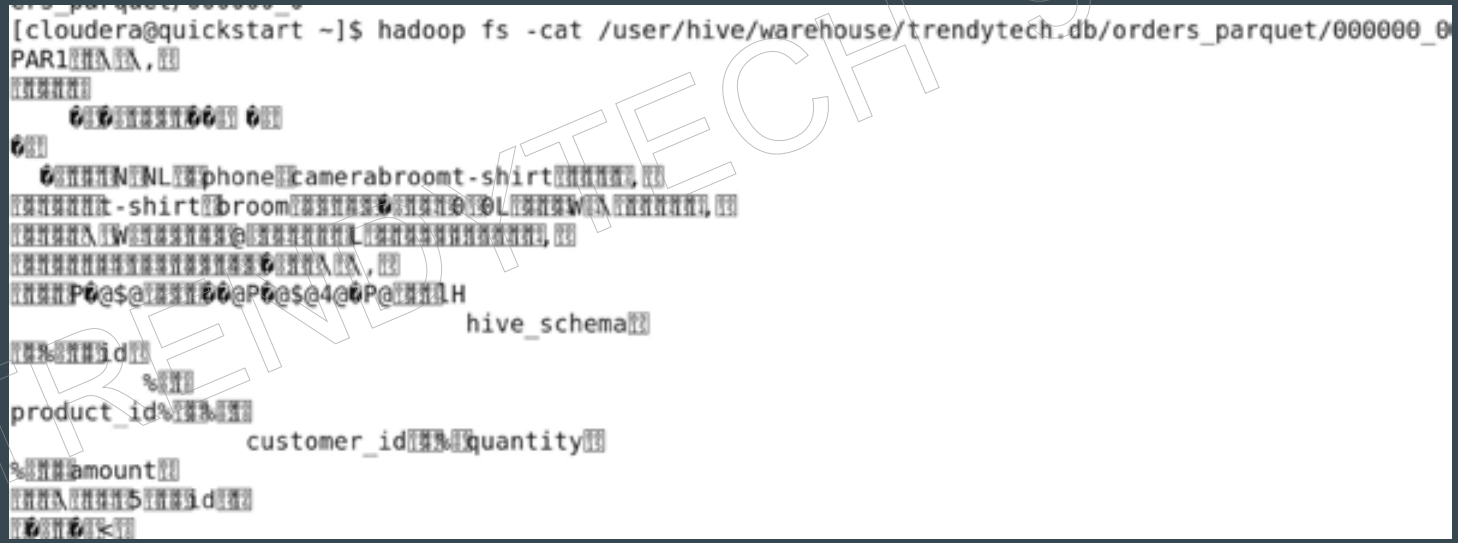

Parquet File Format

Now get the data from hdfs to local using get command.

hadoop fs -get /user/hive/warehouse/trendytech.db/orders_parquet/000000_0 .

Now try to see the metadata using below command.

parquet-tools meta 000000_0

```
[cloudera@quickstart ~]$ parquet-tools meta 000000_0

creator:     parquet-mr version 1.5.0-cdh5.13.0 (build ${buildNumber})

file schema: hive_schema
--------------------------------------------------------------------------------
id:          OPTIONAL INT64 R:0 D:1
product_id:  OPTIONAL BINARY O:UTF8 R:0 D:1
customer_id: OPTIONAL INT64 R:0 D:1
quantity:    OPTIONAL INT32 R:0 D:1
amount:      OPTIONAL DOUBLE R:0 D:1

row group 1: RC:5 TS:429
--------------------------------------------------------------------------------
id:           INT64 UNCOMPRESSED DO:0 FPO:4 SZ:87/87/1.00 VC:5 ENC:RLE,PLAIN,BIT_PACKED
product_id:   BINARY UNCOMPRESSED DO:0 FPO:91 SZ:99/99/1.00 VC:5 ENC:RLE,PLAIN_DICTIONARY,BIT_PACKED
customer_id:  INT64 UNCOMPRESSED DO:0 FPO:190 SZ:88/88/1.00 VC:5 ENC:RLE,PLAIN_DICTIONARY,BIT_PACKED
quantity:     INT32 UNCOMPRESSED DO:0 FPO:278 SZ:68/68/1.00 VC:5 ENC:RLE,PLAIN_DICTIONARY,BIT_PACKED
amount:       DOUBLE UNCOMPRESSED DO:0 FPO:346 SZ:87/87/1.00 VC:5 ENC:RLE,PLAIN,BIT_PACKED
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
```

Parquet File Format

Now try to see the data using below command.

parquet-tools cat 000000_0

```
[cloudera@quickstart ~]$ parquet-tools cat 000000_0
id = 111111
product_id = phone
customer_id = 1111
quantity = 3
amount = 1200.0

id = 111112
product_id = camera
customer_id = 1111
quantity = 1
amount = 5200.0

id = 111113
product_id = broom
customer_id = 1111
quantity = 1
amount = 10.0
```

Json Serde

Ser + de

Serialization + Deserialization

There is no support for Json by default. So we need to add a jar and add that in hive.

Download the jar.

www.congiu.net/hive-json-serde/1.3.7/cdh5/json-serde-1.3.7-jar-with-dependencies.jar

Json Serde

Json Serde

Now add the jar in hive using the below command as show.

```
hive> add jar /home/cloudera/Downloads/json-serde-1.3.7-jar-with-dependencies.jar;
Added [/home/cloudera/Downloads/json-serde-1.3.7-jar-with-dependencies.jar] to class path
Added resources: [/home/cloudera/Downloads/json-serde-1.3.7-jar-with-dependencies.jar]
hive>
```

Json Serde

Create a table using json serde

```
CREATE TABLE orders_json(
  id bigint,
  product_id string,
  customer_id bigint,
  quantity int,
  amount double) ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe';
```

```
hive>  CREATE TABLE orders_json(
    >   id bigint,
    >   product_id string,
    >   customer_id bigint,
    >   quantity int,
    >   amount double) ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe';
OK
Time taken: 0.24 seconds
hive>
```

Json Serde

Now insert the data in this table from orders table

insert overwrite table orders_json select * from orders;

```
hive> insert overwrite table orders_json select * from orders;
Query ID = cloudera_20200507044141_4adaa906-a1bc-4620-9b5e-b1c98f19d44f
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1588847091075_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/applicati
on_1588847091075_0003/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1588847091075_0003
```

Json Serde

Now try to see the data in hdfs

hadoop fs -cat /user/hive/warehouse/trendytech.db/orders_json/*

```
[cloudera@quickstart ~]$ hadoop fs -cat /user/hive/warehouse/trendytech.db/orders_json/*
{"amount":1200,"id":111111,"product_id":"phone","quantity":3,"customer_id":1111}
{"amount":5200,"id":111112,"product_id":"camera","quantity":1,"customer_id":1111}
{"amount":10,"id":111113,"product_id":"broom","quantity":1,"customer_id":1111}
{"amount":20,"id":111114,"product_id":"broom","quantity":2,"customer_id":2222}
{"amount":66,"id":111115,"product_id":"t-shirt","quantity":2,"customer_id":4444}
[cloudera@quickstart ~]$
```

Json Serde

To see all details related to table run the below command

show create table orders_json;

```
CREATE TABLE `orders_json`(
  `id` bigint COMMENT 'from deserializer',
  `product_id` string COMMENT 'from deserializer',
  `customer_id` bigint COMMENT 'from deserializer',
  `quantity` int COMMENT 'from deserializer',
  `amount` double COMMENT 'from deserializer')
ROW FORMAT SERDE
  'org.openx.data.jsonserde.JsonSerDe'
STORED AS INPUTFORMAT
  'org.apache.hadoop.mapred.TextInputFormat'
OUTPUTFORMAT
  'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'
LOCATION
  'hdfs://quickstart.cloudera:8020/user/hive/warehouse/trendytech.db/orders_json'
TBLPROPERTIES (
  'COLUMN_STATS_ACCURATE'='true',
  'numFiles'='1',
  'numRows'='5',
  'rawDataSize'='0',
  'totalSize'='402',
  'transient_lastDdlTime'='1588851707')
Time taken: 0.043 seconds, Fetched: 21 row(s)
```

We have seen hadoop file formats practically

Happy Learning!!!

Follow US

Trainer	Mr. Sumit Mittal
Phone	9108179578
Email	trendytech.sumit@gmail.com
Website	https://trendytech.in/courses/big-data-online-training/
LinkedIn	https://www.linkedin.com/in/bigdatabysumit/
Twitter	@BigdataBySumit
Instagram	bigdatabysumit
Facebook	https://www.facebook.com/trendytech.in/
Youtube	https://www.youtube.com/channel/UCbTggJVf0NDTfWX-C_gUGSg